

# DATA SCIENCE

# REPORT

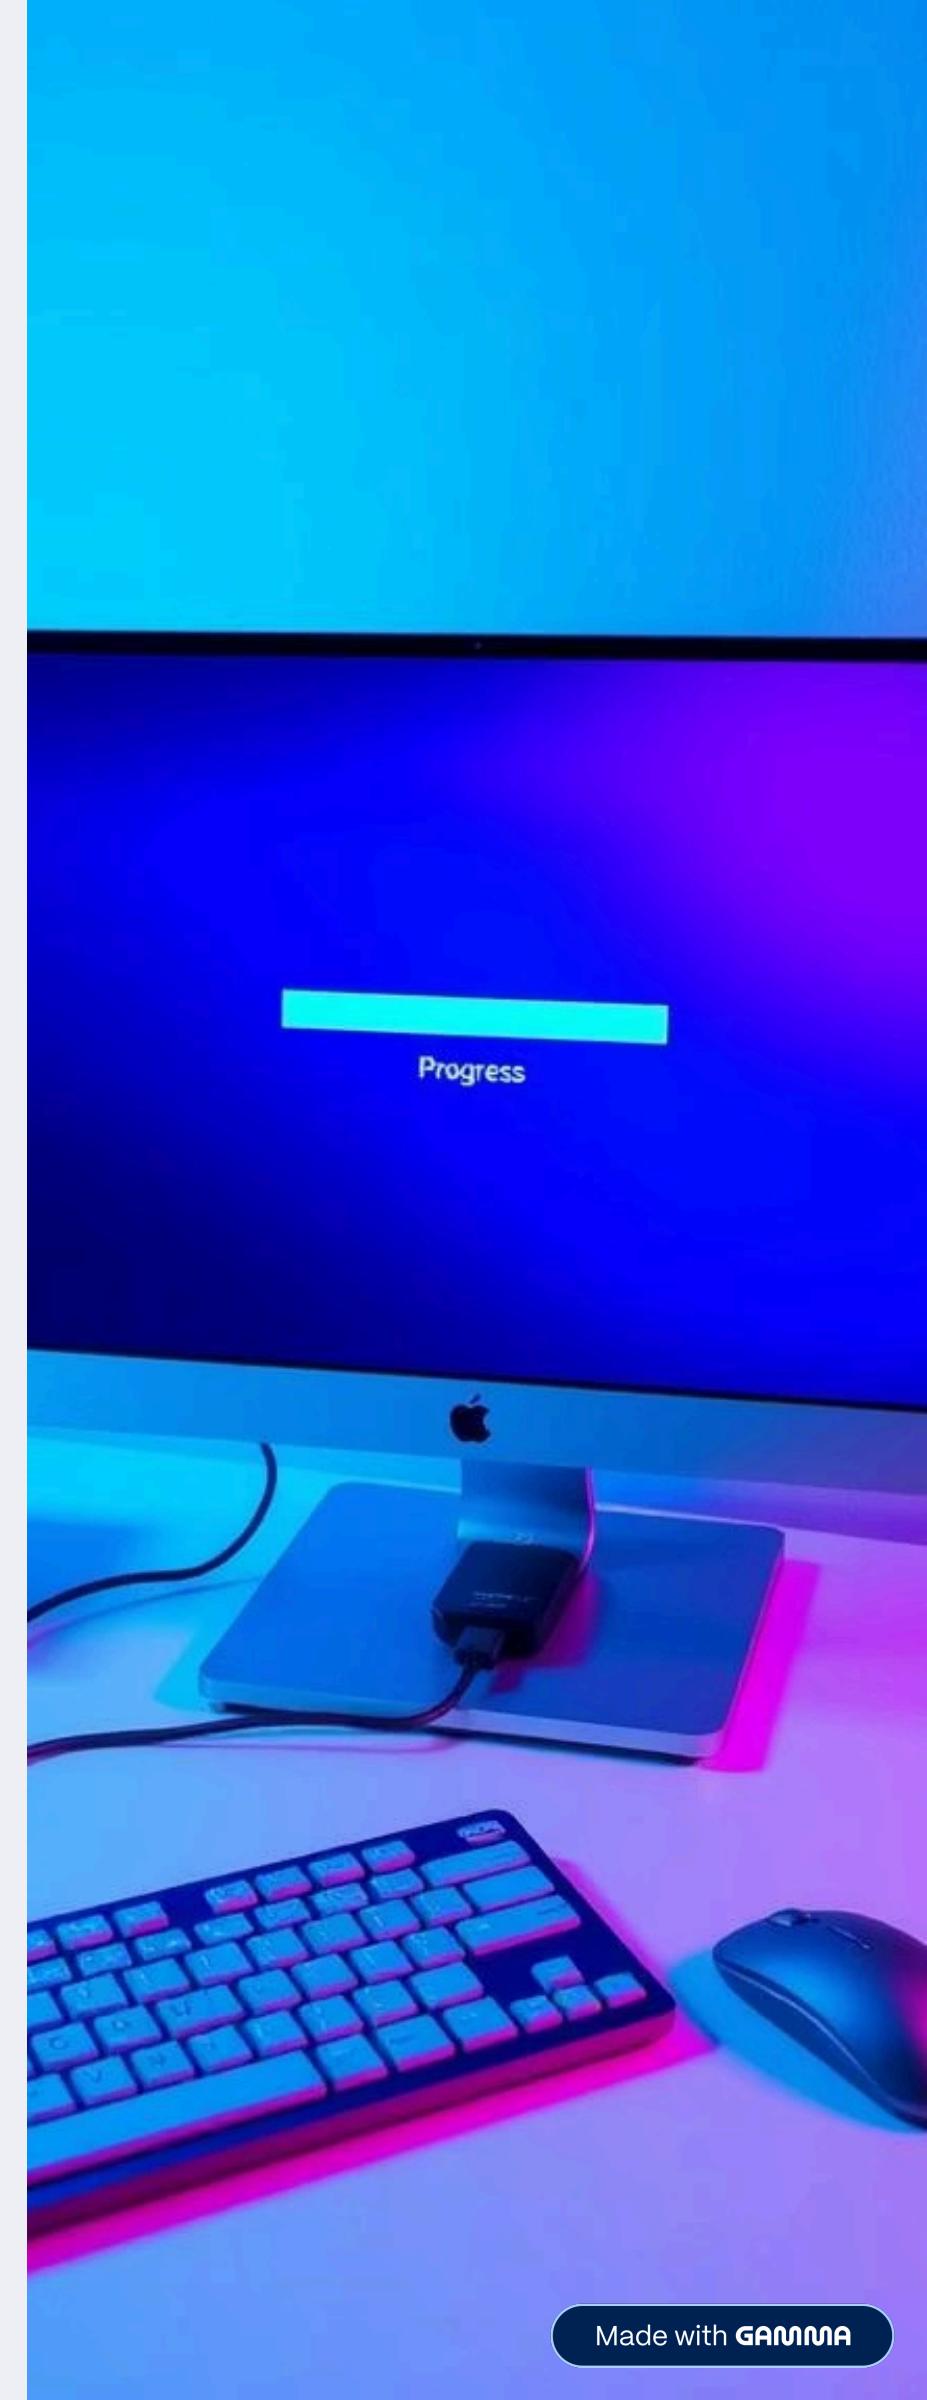
PRESENTED BY :

**Yazan Aqtash**

PRESENTED TO :

**DR.MOHAMMAD AL-HAMMOURI**

**Hashemite University**



# Introduction

**This project aims to prepare a road-accident dataset for machine learning by performing full data cleaning, preprocessing, and target selection.**

**Phase 1 focuses on transforming the raw CSV data into a structured, reliable dataset that can be used for predictive analytics in Phase 2.**

**Using Python (Pandas + Scikit-Learn), several structured steps were performed to ensure the dataset is:**

- Accurate
- Free of duplicates
- Correctly typed
- Enhanced with engineered features
- Ready for machine learning in Phase 2

# Dataset Overview

- **Total rows (before cleaning):** 73,095
- **Total rows (after cleaning):** 72,461
- **Total columns:** 20

## Main Categories of Variables

- **Accident details:**  
Accident Type, Date, Speed (Km), Number of Cars
- **Injury outcomes:**  
Simple Injuries, Medium Injuries, Severe Injuries, Death
- **Environmental & Road conditions:**  
Weather, Light, Road Type, Road Surface Description, Road Properties, Road Lanes
- **Driver & Vehicle info:**  
Driver Age, Driver Licensee Type, Driver Mistake, Sex, Vehicle type, Vehicle Country

## 1. LOAD DATA

### Purpose

The goal of this step is to load the dataset into memory and preserve a raw copy for comparison.

### Code Used

```
df = pd.read_csv("Accident.csv")
df_raw = df.copy()
```

### Explanation

- `df` contains the working dataset.
- `df_raw` is an untouched backup, used to examine data *before* cleaning

## 2. BEFORE-CLEANING SUMMARY

### Purpose

Before cleaning, it is necessary to:

- Check dataset size
- Identify duplicate rows
- Check for missing values
- Understand whether certain columns (e.g., Driver Age) contain unrealistic values

### Code Used

```
print("== BEFORE CLEANING ==")
print(f"Rows: {len(df_raw)}")
print(f"Duplicate rows: {df_raw.duplicated().sum()}")

print("\nMissing values per column:")
print(df_raw.isna().sum())

if "Driver Age" in df_raw.columns:
    print("\n'Driver Age' stats (raw):")
    print(df_raw["Driver Age"].describe())
```

### Explanation

- `df_raw.duplicated().sum()` identifies how many duplicate rows exist.
- `df_raw.isna().sum()` checks for missing values in each column.
- The statistical description of `"Driver Age"` helps detect anomalies such as ages below 10 or above 100.

# 3. DATA CLEANING

The cleaning workflow consists of several essential operations.

## 3.1 Removing Duplicate Rows

### Code Used

```
df = df.drop_duplicates()
```

### Explanation

Duplicate rows distort model learning by overrepresenting certain patterns.  
This step ensures the dataset contains unique entries only.

---

## 3.2 Date Processing

### Code Used

```
df["Date"] = pd.to_datetime(df["Date"], errors="coerce")
df["Year"] = df["Date"].dt.year
df["Month"] = df["Date"].dt.month
df["DayOfWeek"] = df["Date"].dt.dayofweek
df["Hour"] = df["Date"].dt.hour
```

### Explanation

- Converts all date formats (e.g., 14-Jan-18, 2018-01-14) into a standardized `datetime` format.
- Invalid dates become `NaT` due to `errors="coerce"`.
- Extracted features (`Year`, `Month`, `DayOfWeek`, `Hour`) provide meaningful temporal information for modeling:
  - **Year**
  - **Month**: seasonal trends
  - **DayOfWeek**: weekday vs weekend accidents
  - **Hour**: rush hour patterns

## 3.3 Cleaning Driver Age

### Step 1: Identify invalid ages

#### code used

```
df["Driver Age Clean"] = df["Driver Age"].where(  
    (df["Driver Age"] >= 16) & (df["Driver Age"] <= 90)  
)
```

- Ages **below 16 or above 90** were marked invalid
- These values were replaced with `Nan`

### Step 2: Impute missing ages

#### code used

```
median_age = df["Driver Age Clean"].median()  
df["Driver Age Clean"] = df["Driver Age Clean"].fillna(median_age)
```

- Missing/invalid ages were filled using the **median age (34)**
- This preserves the natural distribution without distortion

---

## 3.4 Creating the Target Variable

#### Code Used

```
df["Injury_or_Death"] = (  
    (df["Simple Injuries"] > 0) |  
    (df["Medium Injuries"] > 0) |  
    (df["Severe Injuries"] > 0) |  
    (df["Death"] > 0)  
)astype(int)
```

#### Explanation

The dataset originally contained **exact counts of injuries and deaths**.

This step converts those multiple outcome columns into a **single binary classification target**:

- **1 = Accident resulted in at least one injury or death**
- **0 = Accident resulted in no injuries and no deaths**

This simplifies the predictive modeling process.

## 3.5 Dropping Leakage and Replaced Columns

### Code Used

```
drop_cols = [  
    "Simple Injuries", "Medium Injuries", "Severe Injuries",  
    "Death", "Date", "Driver Age"  
]  
df = df.drop(columns=drop_cols)
```

### Explanation

Removed columns:

- Simple Injuries
- Medium Injuries
- Severe Injuries
- Death
- Date (replaced by Year/Month/etc.)
- Driver Age (replaced with Driver Age Clean)

Columns like "Simple Injuries" or "Severe Injuries" occur *after* the accident and would create **data leakage** if used for prediction.

## 4. AFTER-CLEANING SUMMARY

### Code Used

```
print("\n==== AFTER CLEANING ====")
print(f"Rows: {len(df)}")

print("\nMissing values per column:")
print(df.isna().sum())

print("\n'Driver Age Clean' stats (cleaned):")
print(df["Driver Age Clean"].describe())

print("\nTarget 'Injury_or_Death' distribution:")
print(df["Injury_or_Death"].value_counts(normalize=True).rename("proportion"))
```

### Explanation

This step verifies:

- Final dataset size
- Remaining missing data (should be none)
- Cleaned age distribution
- How imbalanced the `Injury_or_Death` target is (fatal accidents are very rare)

## 5. SAVE CLEANED DATA

### Code Used

```
df.to_csv("Accident_cleaned.csv", index=False)
```

### Explanation

The cleaned dataset is exported for future use in modeling.

# 6. FEATURE AND TARGET SEPARATION

## Code Used

```
X = df.drop(columns=["Injury_or_Death"])
y = df["Injury_or_Death"]
```

## Explanation

- **X** contains all predictor (feature) variables
- **y** contains the binary target variable (**Injury\_or\_Death**)
- The target equals **1** if an accident resulted in **any injury or death**, and **0** otherwise
- This formulation avoids extreme class imbalance caused by rare fatal accidents and enables meaningful model learning

**Due to the very low number of fatal accidents, injury and death outcomes were merged into a single severity-based target variable.**

## 6.1 Identifying Numerical and Categorical Columns

## Code Used

```
num_cols = X.select_dtypes(include=["int64", "float64"]).columns.tolist()
cat_cols = X.select_dtypes(include=["object"]).columns.tolist()
```

## Explanation

Scikit-learn preprocessing requires knowing which columns are numeric vs categorical:

- Numeric → StandardScaler
- Categorical → OneHotEncoder

## 7. PREPROCESSOR PIPELINE

### Code Used

```
preprocess = ColumnTransformer(  
    transformers=[  
        ("num", StandardScaler(), num_cols),  
        ("cat", OneHotEncoder(handle_unknown="ignore"), cat_cols),  
    ]  
)
```

### Explanation

This pipeline does two things:

- **Scale** all numerical features (centers and normalizes values)
- **One-hot encode** all categorical features into numeric binary vectors

Using `handle_unknown='ignore'` ensures unseen categories at prediction time do not break the model.

## 8. TRAIN / TEST SPLIT

### Code Used

```
X_train, X_test, y_train, y_test = train_test_split(  
    X,  
    y,  
    test_size=0.2,  
    random_state=42,  
    stratify=y  
)
```

### Explanation

- Split is **80% training / 20% testing**
- `stratify=y` ensures both sets keep the same proportion of fatal vs non-fatal accidents
- `random_state=42` ensures reproducibility

# 9. OPTIONAL PREPROCESSOR TRANSFORMATION

## Code Used

```
X_train_trans = preprocess.fit_transform(X_train)
```

```
X_test_trans = preprocess.transform(X_test)
```

## Explanation

If uncommented, this would transform training and testing data according to the preprocessing pipeline. This is typically done right before training machine learning models in Phase 2

## 2. BEFORE-CLEANING OUTPUT

```
== BEFORE CLEANING ==
```

```
Rows: 73095
```

```
Duplicate rows: 634
```

```
Missing values per column:
```

Accident Type	0
Date	0
Speed (Km)	0
Simple Injuries	0
Severe Injuries	0
Death	0
Medium Injuries	0
Road Lanes	0
Road Surface Description	0
Vehicle Country	0
Driver Licensee Type	0
Road Type	0
Light	0
Weather	0
Road Properties	0
Driver Mistake	0
Vehicle type	0
Number of Cars	0
Driver Age	0
Sex	0

```
dtype: int64
```

```
'Driver Age' stats (raw):
```

count	73095.000000
mean	35.497353
std	23.094740
min	0.000000
25%	24.000000
50%	33.000000
75%	44.000000
max	118.000000

```
Name: Driver Age, dtype: float64
```

# AFTER-CLEANING OUTPUT

```
== AFTER CLEANING ==
```

```
Rows: 72461
```

```
Missing values per column:
```

Accident Type	0
Speed (Km)	0
Road Lanes	0
Road Surface Description	0
Vehicle Country	0
Driver Licensee Type	0
Road Type	0
Light	0
Weather	0
Road Properties	0
Driver Mistake	0
Vehicle type	0
Number of Cars	0
Sex	0
Year	94
Month	94
DayOfWeek	94
Hour	94
Driver Age Clean	0
Injury_or_Death	0
dtype: int64	

```
'Driver Age Clean' stats (cleaned):
```

count	72461.000000
mean	36.132485
std	11.283246
min	18.000000
25%	28.000000
50%	34.000000
75%	42.000000
max	90.000000

```
Name: Driver Age Clean, dtype: float64
```

```
Target 'Injury_or_Death' distribution:
```

Injury_or_Death	
0	0.946233
1	0.053767

```
Name: proportion, dtype: float64
```

```
Cleaned data saved to 'Accident_cleaned.csv'
```

```
Numeric columns: ['Speed (Km)', 'Number of Cars', 'Sex', 'Year', 'Month', 'DayOfWeek', 'Hour', 'Driver Age Clean']
```

```
Categorical columns: ['Accident Type', 'Road Lanes', 'Road Surface Description', 'Vehicle Country', 'Driver Licensee Type', 'Road Type', 'Light', 'Weather']
```

```
== DATA SPLIT SHAPES ==
```

X_train:	(57968, 19)
X_test:	(14493, 19)
y_train:	(57968,)
y_test:	(14493,)

```
Process finished with exit code 0
```

# 10. Conclusion

- Dataset loaded and examined
- Removed duplicates
- Mixed-format dates standardized and decomposed
- Driver ages cleaned and invalid entries corrected
- Binary target variable created (Has\_Death)
- Removed leakage columns
- Classified columns into numerical and categorical
- Scikit-learn preprocessing pipeline constructed
- Train/test split executed
- Cleaned dataset exported